

Attempt all the questions.

[CLO 1: Understand and design the structure of deep neural networks]

Q1: Analyze the given code to identify any errors and provide a corrected version of the code.

[10 marks]

	Code	Answer
1	<pre>import torch.nn as nn import torch class NeuralNetwork(nn.Module): def __init__(self, input_size, hidden_size, num_classes): super(SimpleNet, self).__init__() self.fc1 = nn.Linear(input_size, hidden_size) self.relu = ReLU() self.fc2 = nn.Linear(hidden_size, num_classes) def forward(self, x): x = self.fc1(x) x = self.relu x = self.fc2(x) model = SimpleNet(784, 100, 10) print(model(torch.rand(1, 784)).shape)</pre>	

inputs = inputs.view(inputs.size(0), -1)
outputs = model(inputs)
loss = criterion(outputs, labels)

loss.backward()
optimizer.zero_grad()
optimizer.step()

print(f"Epoch {epoch + 1}, Loss: {loss.item():.4f}")

it have
also zero g
should b
above 0

3. import torch
from torch.utils.data import DataLoader, TensorDataset

x_data = torch.rand(100, 10)
y_data = torch.randint(0, 2, (100,))
dataset = TensorDataset(x_data, y_data)

train_loader = DataLoader(dataset, batch_size
= 16, shuffle = True)

for batch in train_loader:
 inputs = batch[:, :-1]
 labels = batch[:, -1]
 print(inputs.shape, labels.shape)

there is
issue in
y-data

4. import torch
x = torch.rand(32, 3, 28, 28)
x_reshaped = x.view(32, -1, 28)
print(x_reshaped.shape)

2025/05/06 14:57

5.	<pre>import torch a = torch.tensor([1, 2], [3, 4]) b = a.unsqueeze(0) print(a.shape, b.shape)</pre>	
6.	<pre>import torch x = torch.tensor(5.0, requires_grad = True) y = x ** 2 z = y.detach() print(y.requires_grad, z.requires_grad)</pre>	
7.	<pre>import torch x = torch.tensor(4.0, requires_grad = True) y = x ** 2 + 3 * x + 5 y.backward() print(x.grad)</pre>	
8.	<pre>import torch x = torch.tensor(1.0, requires_grad = True) y = x ** 3 + x ** 2 y.backward(retain_graph = True) y.backward() print(x.grad)</pre>	2025/05/06 14:58
9.	<pre>import torch x = torch.tensor(2.0, requires_grad = True) y = x ** 3 for i in range(2): y.backward() print(x.grad)</pre>	
10.	<pre>import torch import numpy as np np_array = np.array([1, 2, 3, 4]) tensor = torch.from_numpy(np_array) tensor[2] = 10</pre>	

National University of Computer and Emerging Sciences Islamabad Campus

Part B: Now, consider the same neural network, but replace the SVM loss with Cross-Entropy Loss. Use the above given resulting logits matrix and the corresponding true class labels to compute the Cross-Entropy loss of each example and calculate the gradient of the loss w.r.t all input logits.

Note: Don't need to show all the steps.

Now Find \rightarrow cross entropy loss
creating an output matrix

33.11	16.47	11.80	60.68
42.77	40.14	10.93	181.27
244.4	897.8	9897.1	12.18
27.1	16.47	812.1	90

\rightarrow Find \rightarrow exp

0	0.02	0.006	0.96	0.039	1540.62
2	0.042	0.003	0.938	0.01	11652.45
1	0.022	0.008	0.895	0.000	11052.89
0	0.02	0.017	0.89	0.095	945.95

10

Find \rightarrow loss

$$-1 \log(0.02) + 1 \log(0.9386) + 1 \log(0.081) + 1 \log(0.02)$$

$$= (1.698 + 0.02 + 1.09 + 2.698)$$

calculated 155.50697

Find \rightarrow gradients

-0.98	0.006	0.96	0.039
0.042	0.039	-0.063	0.01
0.022	-0.11	0.895	0.001

	f_1	f_2	f_3	f_4	Y
1.	<u>3.5</u>	2.8	7.3	4.2	0
2.	6.2	3.7	<u>9.3</u>	5.2	2
3.	5.5	6.8	9.2	2.5	1
4.	<u>3.3</u>	2.8	6.7	4.5	0

In case of multiclass sum loss

$$\max(0, 2.8 - 3.5), \max(0, 7.3 - 3.5), \max(0, 4.2 - 3.5)$$

0	0	3.8	0.7
0	0	0	0
0	0	2.4	0
0	0	3.4	1.2

$$\text{Total loss} = 11.5$$

logits matrix

No to finding derivatives

-2	0	1	1
0	0	0	0
0	-1	1	0
-2		1	1

10

[CLO 2: Understand the different layers and their operation]

Q3: Use the space provided for answering the short questions. Be precise while answering the short questions.

[10 marks]

(a) Prove the SoftMax loss contribution from a single example $L_i = -\log \frac{e^{f_i}}{\sum_j e^{f_j}}$ can be equivalently written as $L_i = -e^{f_i} + \log \sum_j e^{f_j}$.

in case of $\log \log_a$ can be written as

$$-\log \frac{e^{f_i}}{\sum_j e^{f_j}}$$

$$= -(\log e^{f_i} - \log \sum_j e^{f_j})$$

$$= -(\log e) \cdot e^{f_i} - \log \sum_j e^{f_j}$$

$$= -e^{f_i} + \log \sum_j e^{f_j}$$

(b) Difference between gradient descent, stochastic gradient descent, and batch gradient descent, which is more preferred one and why?

normal gradient descent calculate ~~outputs~~
~~from backward function~~ derivatives in back
propagation

(c) Why do we need derivatives w.r.t to inputs while computing back-propagation for affine layer.


because to compute value in backward pass
we multiply all gradients local derivatives
that's why we calculate gradients

**National University of Computer and Emerging Sciences
Islamabad Campus**

(d) How does the loss function of a Multi-Class SVM differ from Cross-Entropy Loss?

In case of multi class svm loss we find difference of each element ~~than given~~ with y index. ~~given~~ compare it with zero. In case if it is greater than zero we consider that value and we sum up at the end in case of CSE Loss we find exponentials and then sum up and find probability and of y labels and sum them up at last.

(e) Why is the softmax function typically used before applying Cross-Entropy Loss in multi-class classification?




[CLO 2: Understand the different layers and their operation]

Q4: [30 marks]

Part A: Consider a 3-layer neural network with a multi-class SVM loss function (architecture diagram is given in Figure 1). The network's final layer consists of 4 neurons, each representing a different output class. We input 4 examples into the network, and each neuron produces a logit score for each example. The resulting logits matrix and the corresponding true class labels are provided below. Compute the multi-class SVM loss of each example and calculate the gradient of the loss w.r.t all input logits.

Note: Don't need to show all the steps.



part C: Now, consider the same neural network, but replace the SVM loss with Cross-Entropy Loss. Implement a function to compute the loss and its gradient based on the given logits and actual labels.

```
def cross_entropy_loss(logits, y):
```

"""

Computes the cross - entropy loss and its gradient.

Parameters:

logits (Tensor): Predicted logits from the model.

y (Tensor): Ground - truth class labels.

Returns:

loss (float): Scalar value representing the cross - entropy loss.

grad (Tensor): Gradient of the loss with respect to the input logits.

"""

critereon = torch.nn.CrossEntropyLoss()

def cross-Entropy loss (logits, y)

exp-matrix = np.exp (logits)

probabilities-matrix = ~~logits~~ / np.sum (exp-matrix, axis=1)

probability-matrix = ~~logits~~ / probability-matrix

loss = -log (np.sum (probability-matrix * np.exp(logits), y))

logit-matrix = ~~logits~~

return ~~loss~~

method 2

alternative method using library

critereon = torch.nn.CrossEntropyLoss()

def cross-Entropy loss (logits, y)

loss = critereon (logits, y)

return loss